

Feature Pyramid Networks for Object Detection (2017)

Tsung-Yi Lin • Piotr Dollar • Ross Girshick
Kaiming He • Bharath Hariharan • Serge Belongie

Facebook AI Research
Cornell University and Cornell Tech

Presented by
Selahaddin HONİ



Digital Video Processing
Spring 2021

Content

Introduction

Feature Maps

Image Pyramid & Sliding Windows

Literature Review

Featurized Image Pyramid :: DPM

Single Feature Map :: Faster R-CNN

Pyramidal Feature Hierarchy :: SSD

Feature Pyramid Networks



Theory

Simple Illustration

Applications: RPN


Applications: Fast R-CNN

Benchmark




Semantic value increases with higher level features


Image pyramid is simply the scale variations of an image. Main purpose is detecting an object in different scales. Sliding windows is a basic search mechanism.




Proposed network architectures for detecting in different scales




(a) Featurized image pyramid




(b) Single feature map



(c) Pyramidal feature hierarchy



(d) Feature Pyramid Network



Problem is too slow.

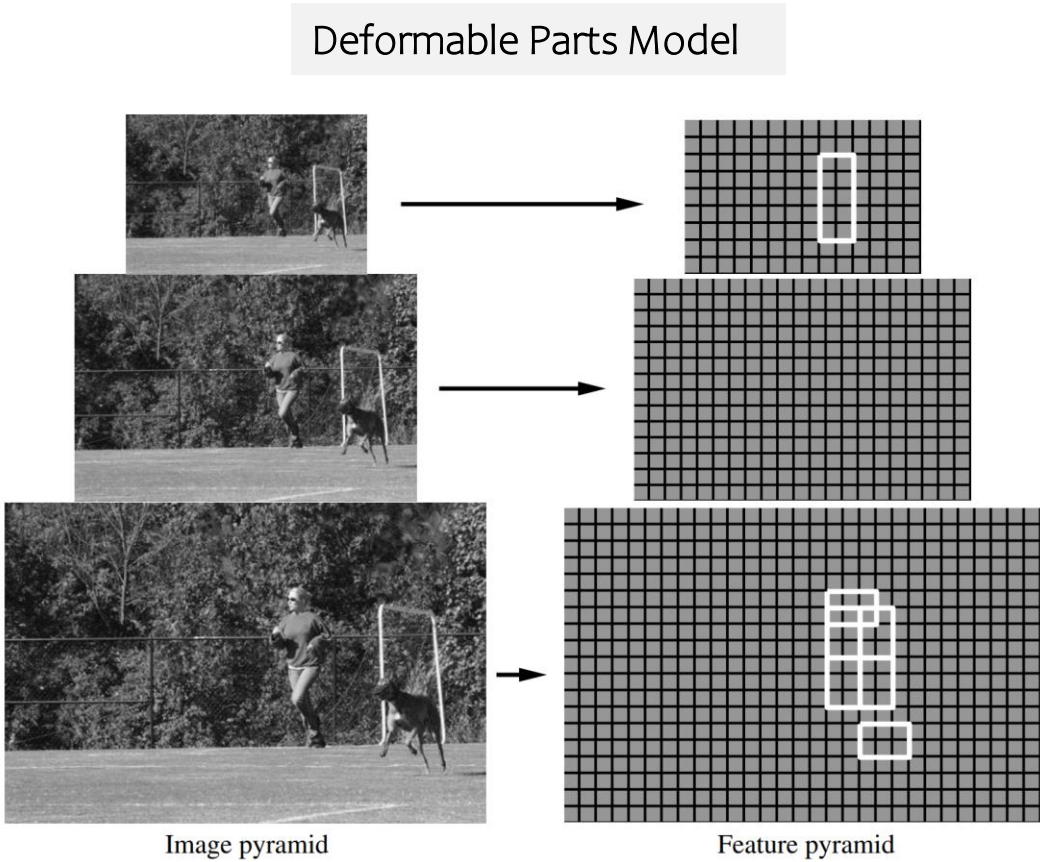




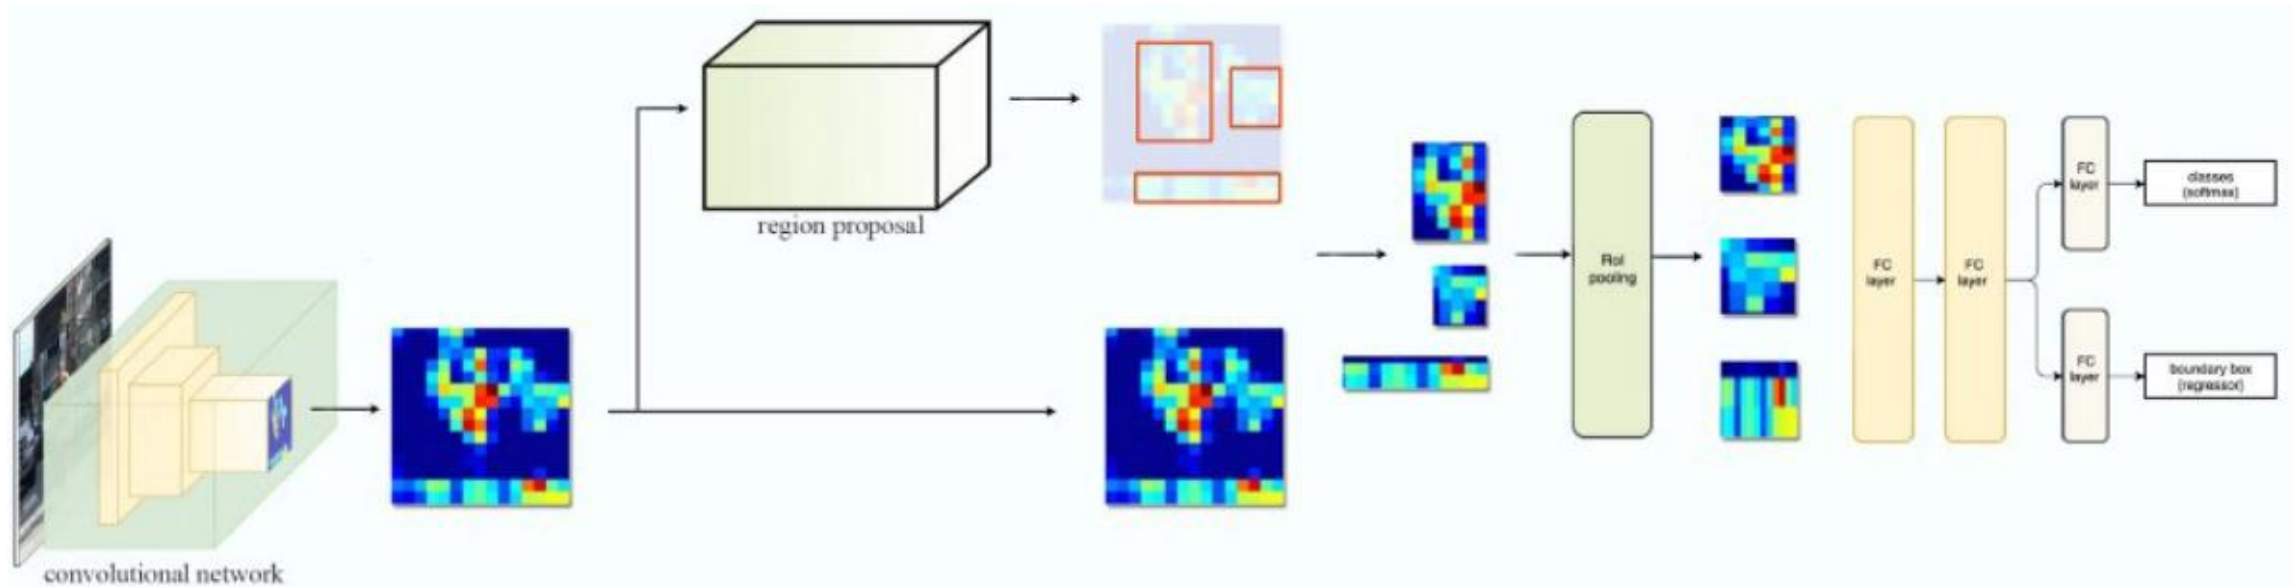
Fig. 3. A feature pyramid and an instantiation of a person model within that pyramid. The part filters are placed at twice the spatial resolution of the placement of the root.



Problem

RPN only use the final feature map of the backbone to propose the regions.


Faster R-CNN




Input Image




First Conv Block (512x512)




Second Conv Block (256x256)

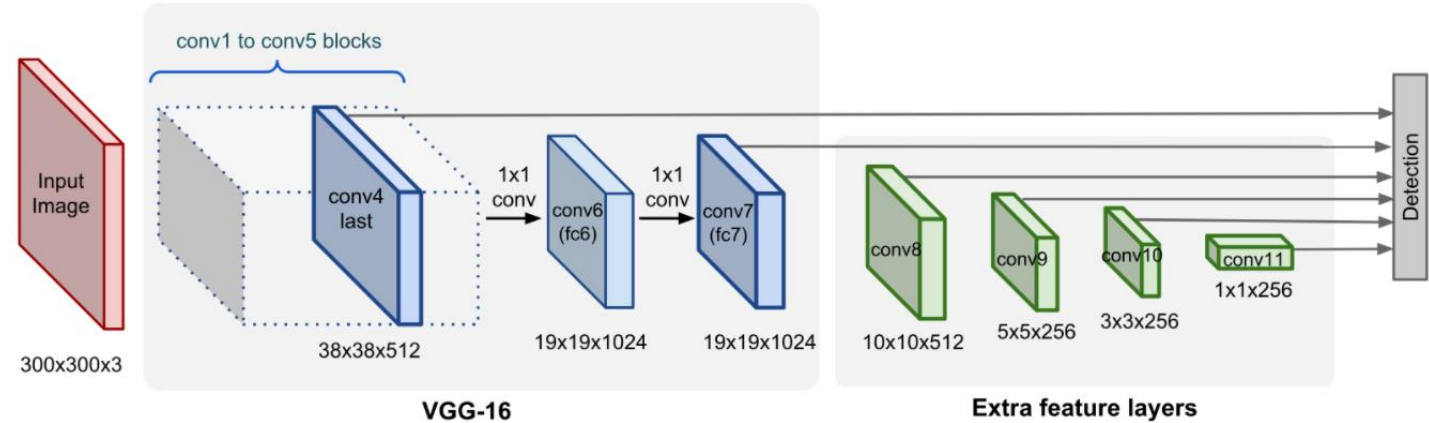


Third Conv Block (128x128)





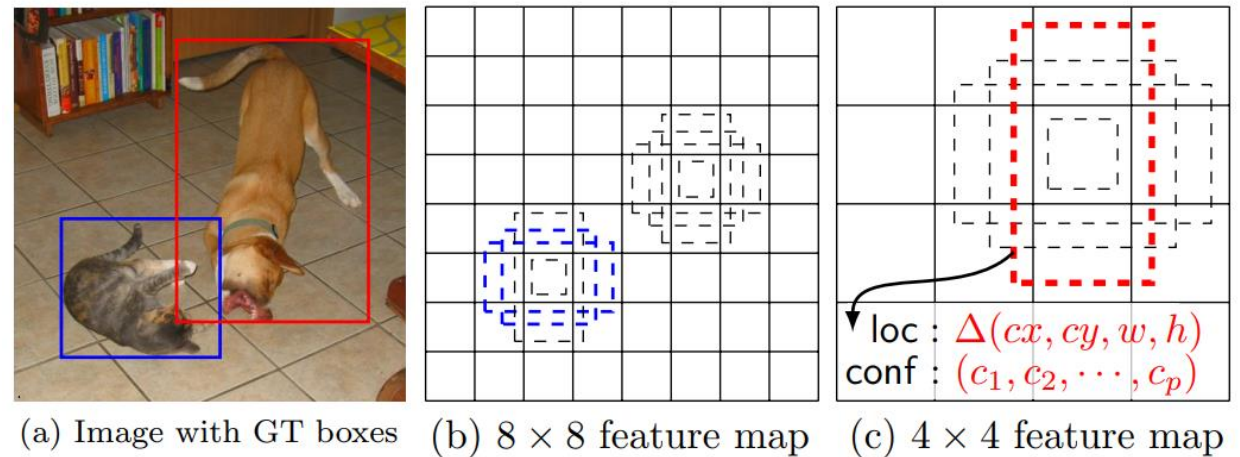
Single Shot Detector



SSD predicts offset of predefined anchor boxes ('default boxes' in the paper[1]) for every location of the feature map.

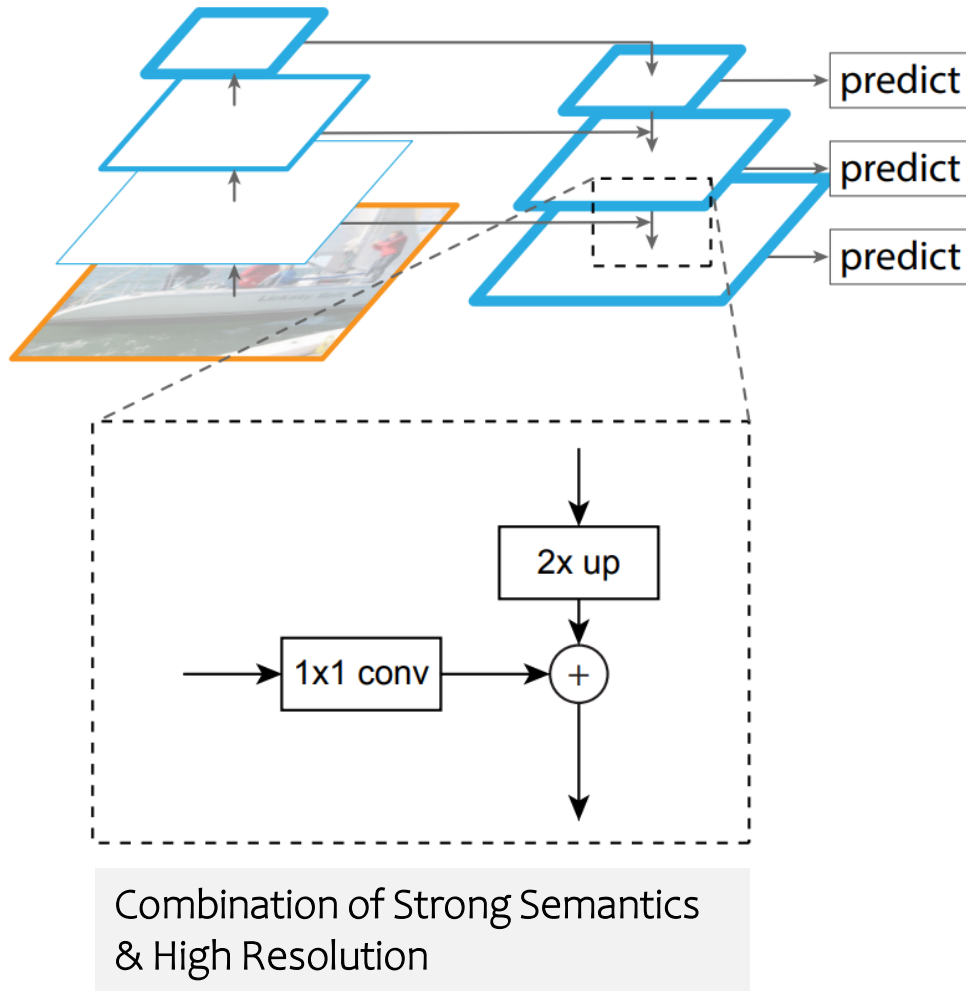
Problem is SSD misses to use the low-level features or high resolution map. It is important for detecting small objects.

Hierarchy starts from at the end of the backbone.




[1] Wei Liu, Dragomir Anguelov, Dumitru Erhan, Christian Szegedy, Scott Reed, Cheng-Yang Fu, Alexander C. Berg (2016). "SSD: Single Shot MultiBox Detector"


[2] [Object Detection Part 4: Fast Detection Models \(lilianweng.github.io\)](https://lilianweng.github.io/)



Upsampling (Nearest Neighbor)

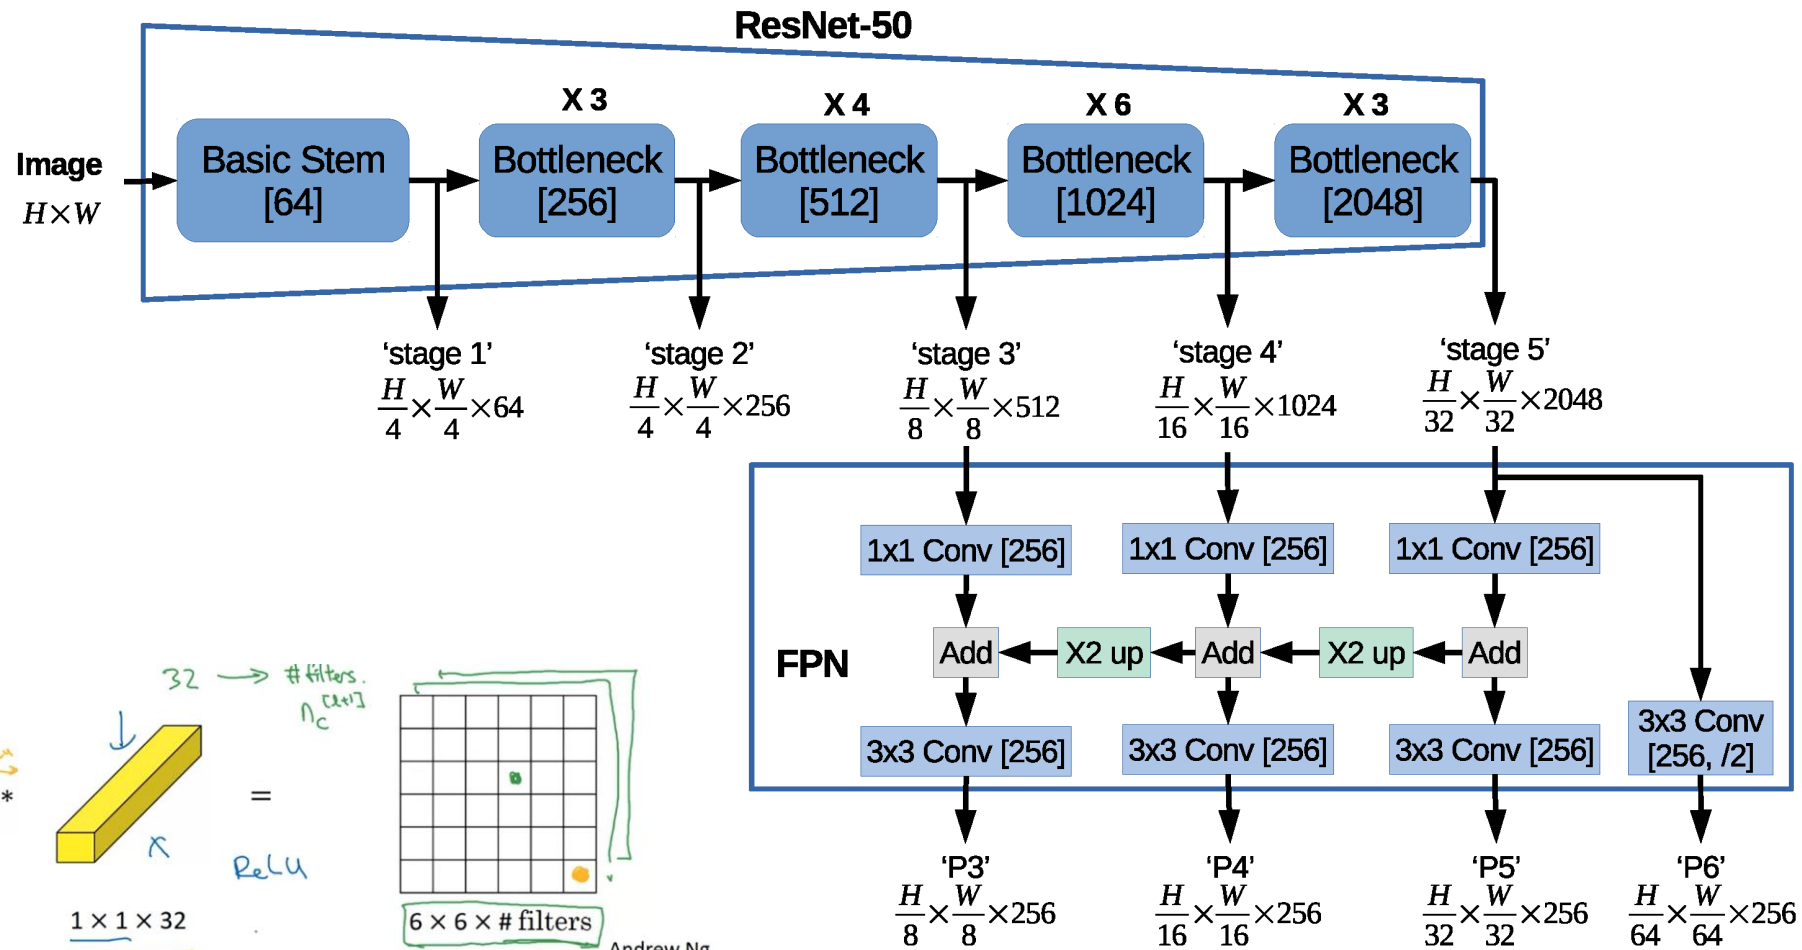


1x1 conv adjusts the channel dimension and adds non-linearity

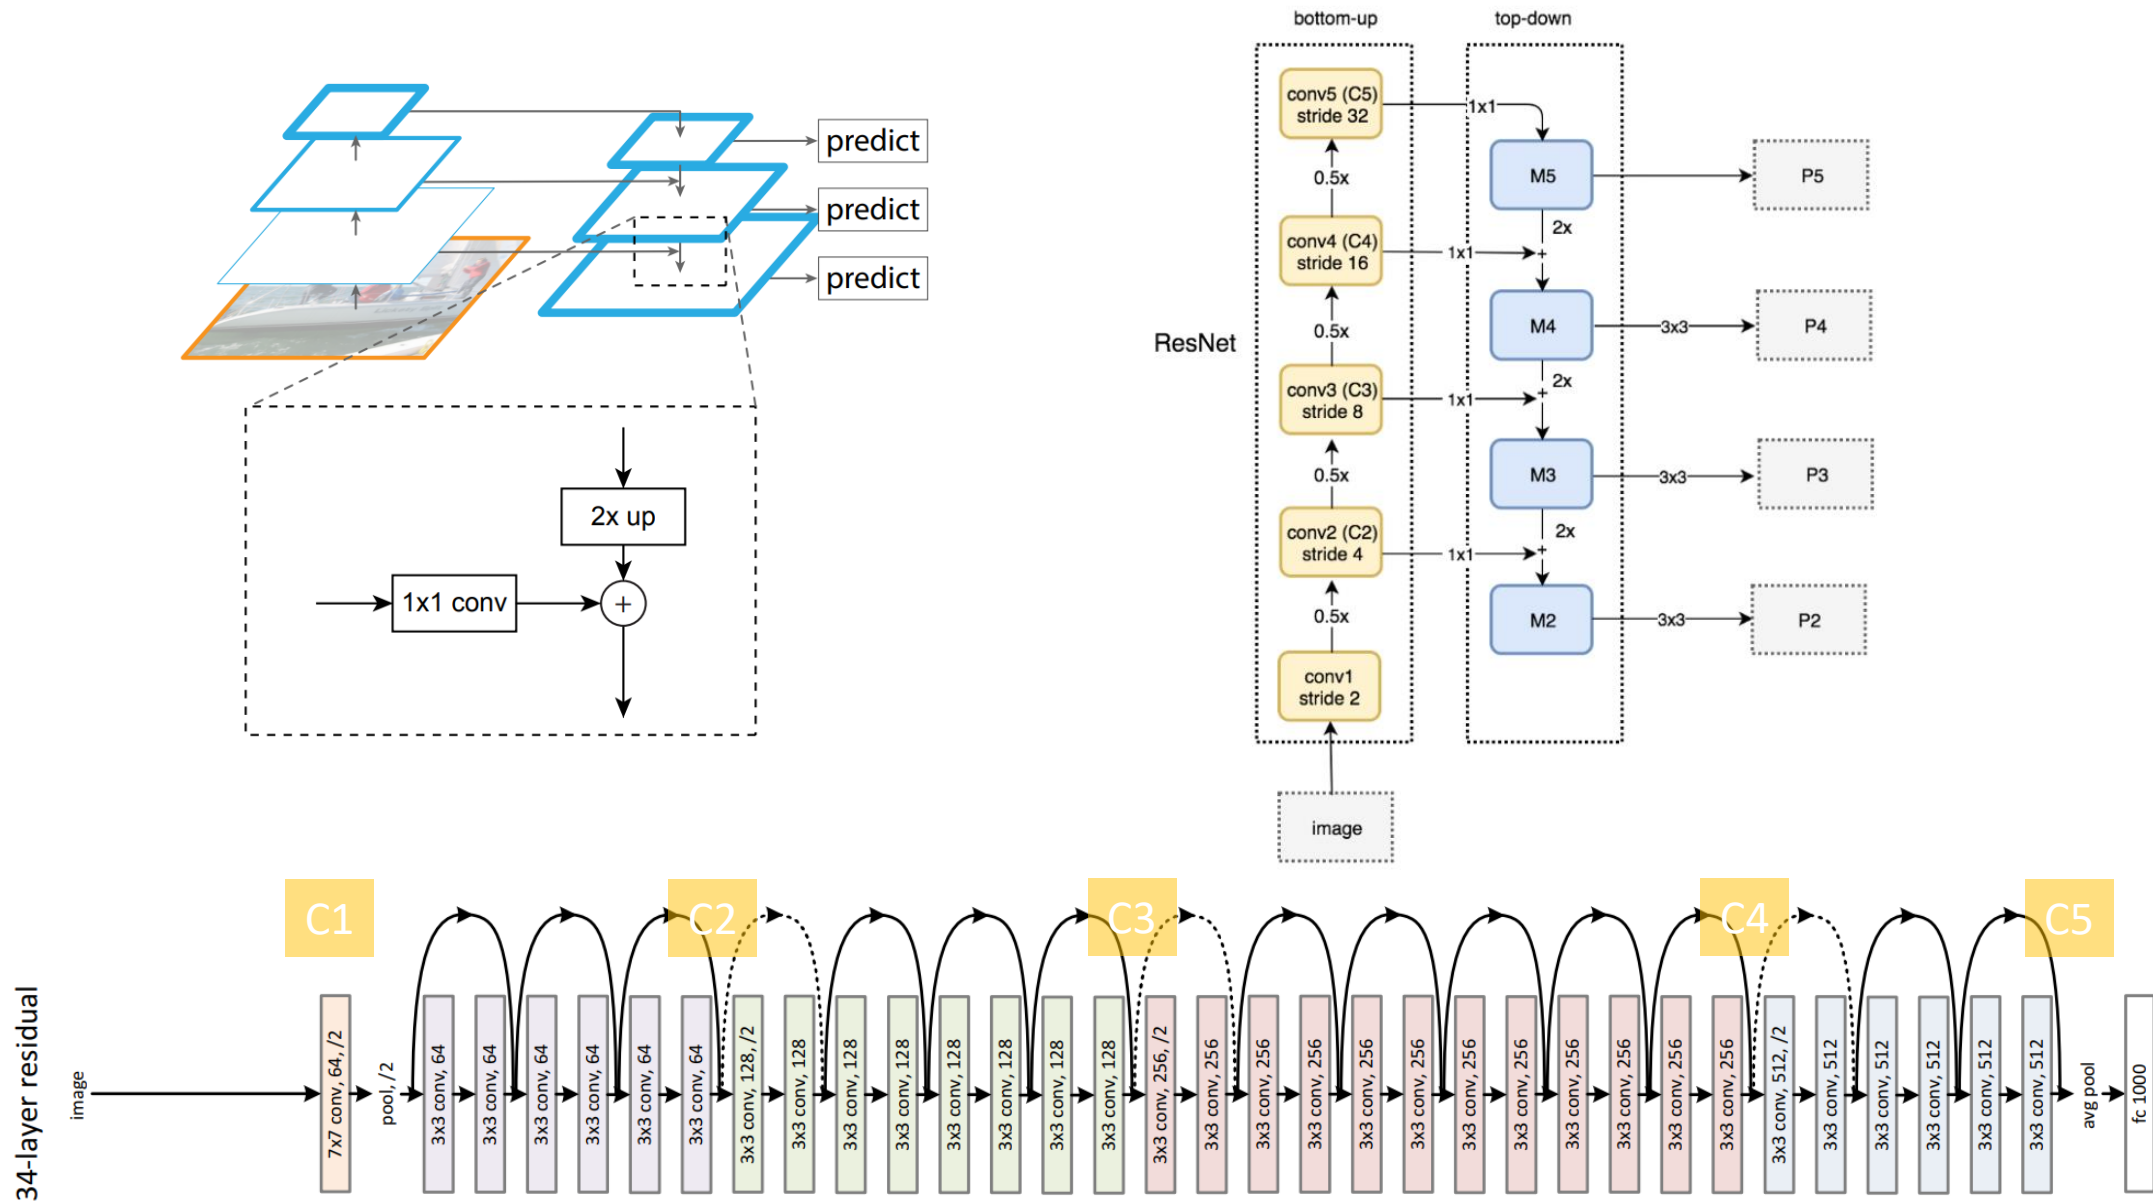


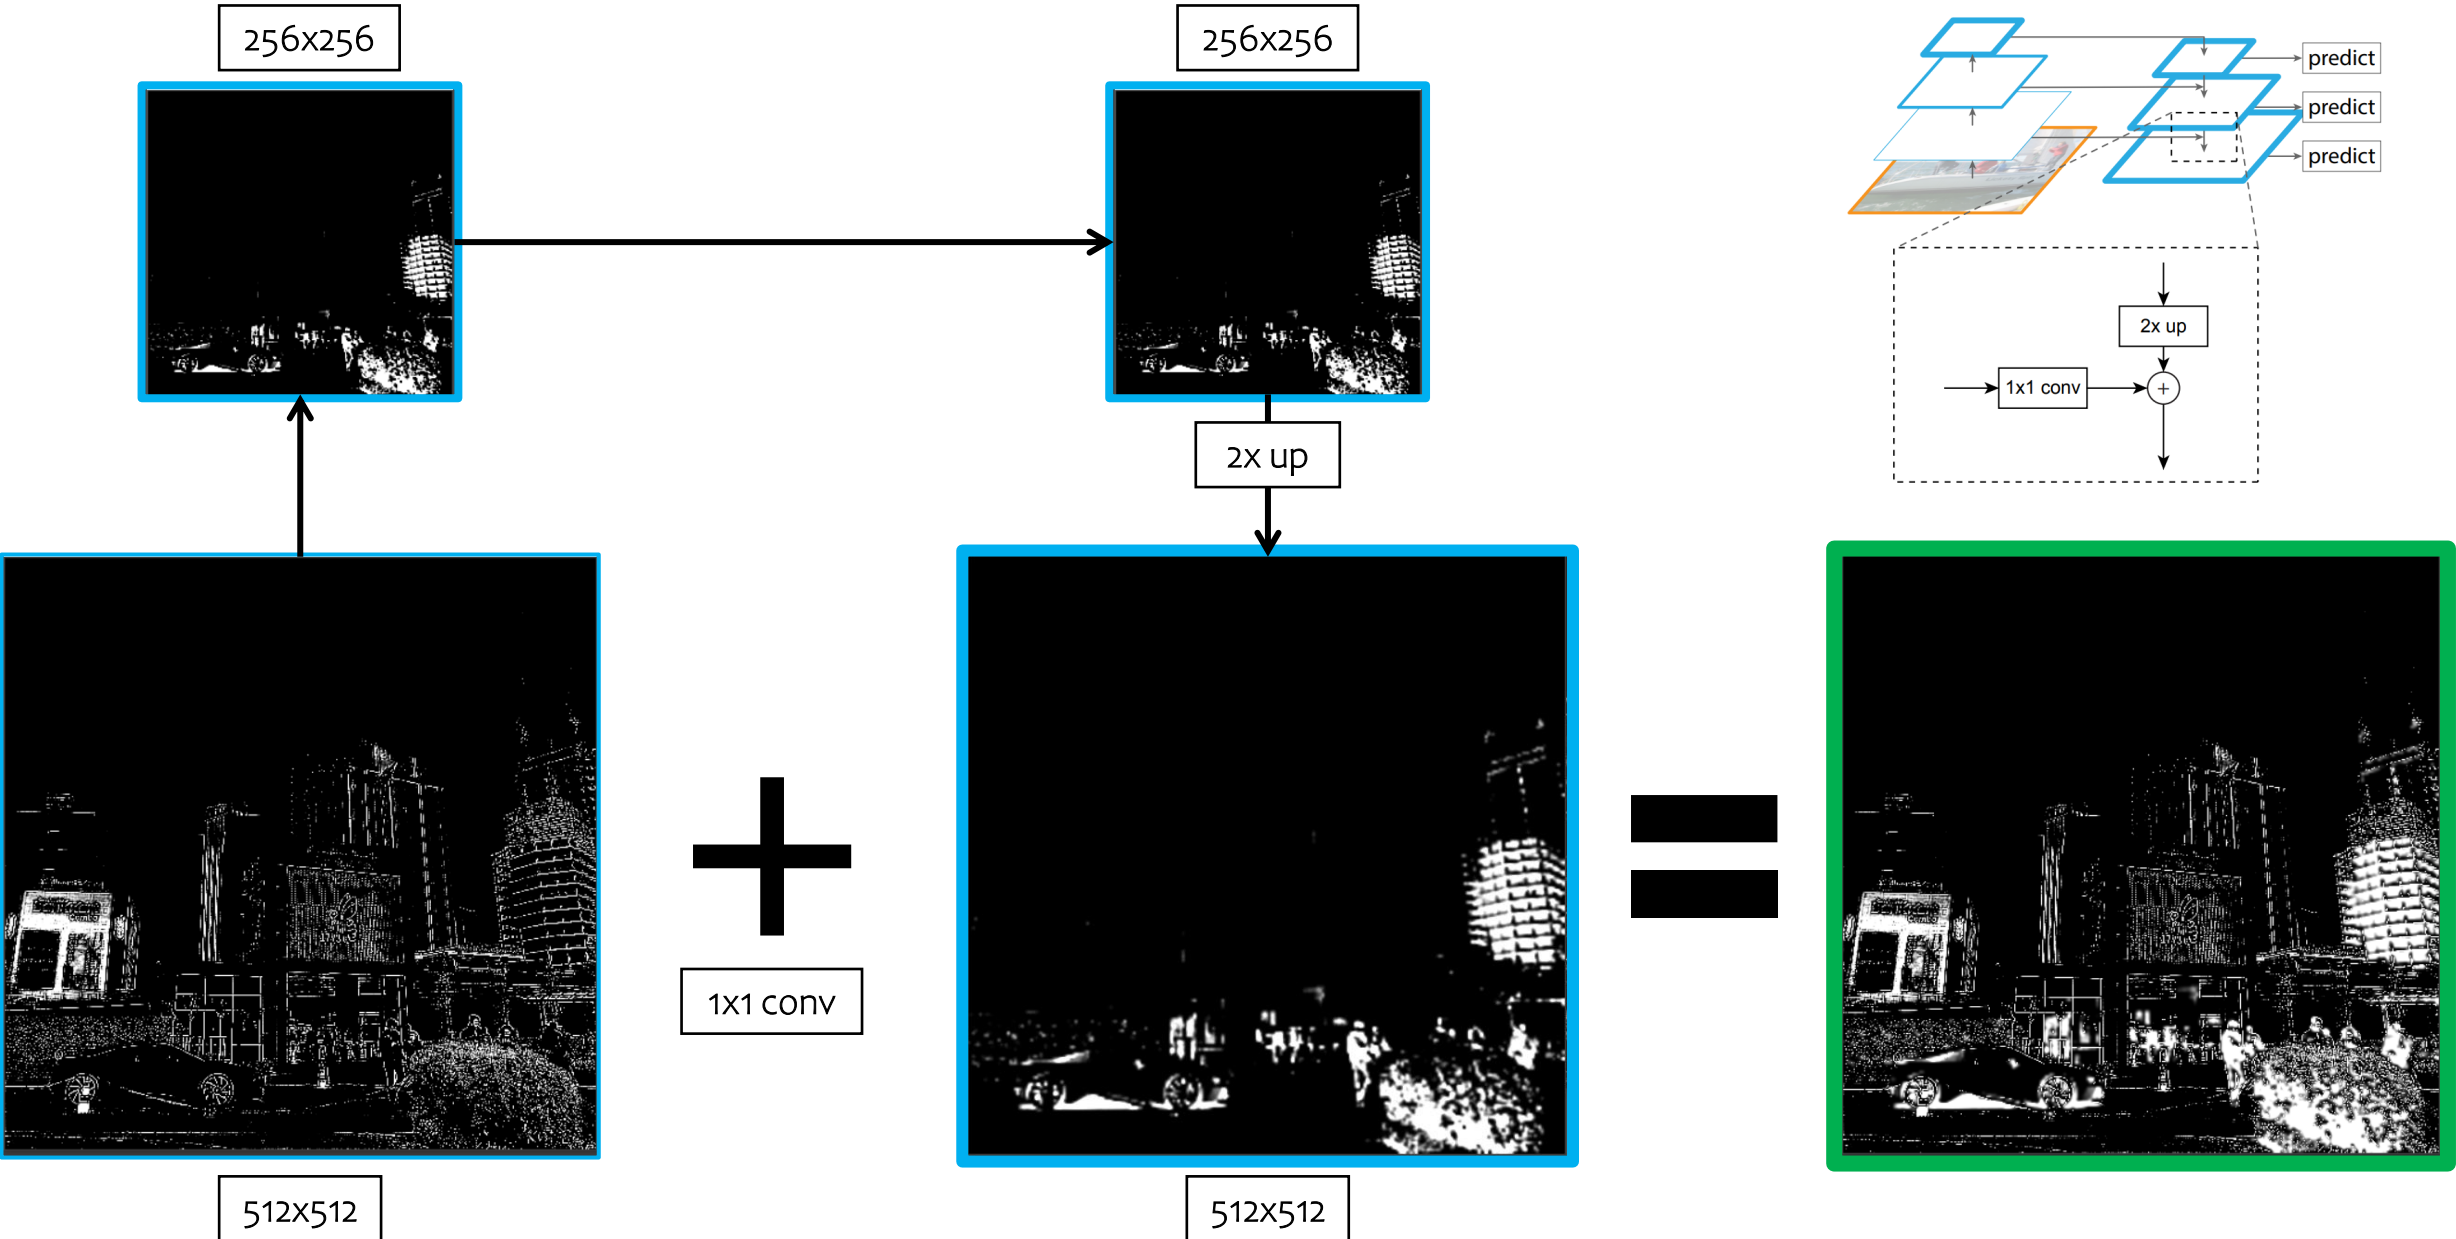
[1] [YasinEnigma/Image Interpolation: Image interpolation implementation using pure python and compare result with Opencv. \(github.com\)](#)

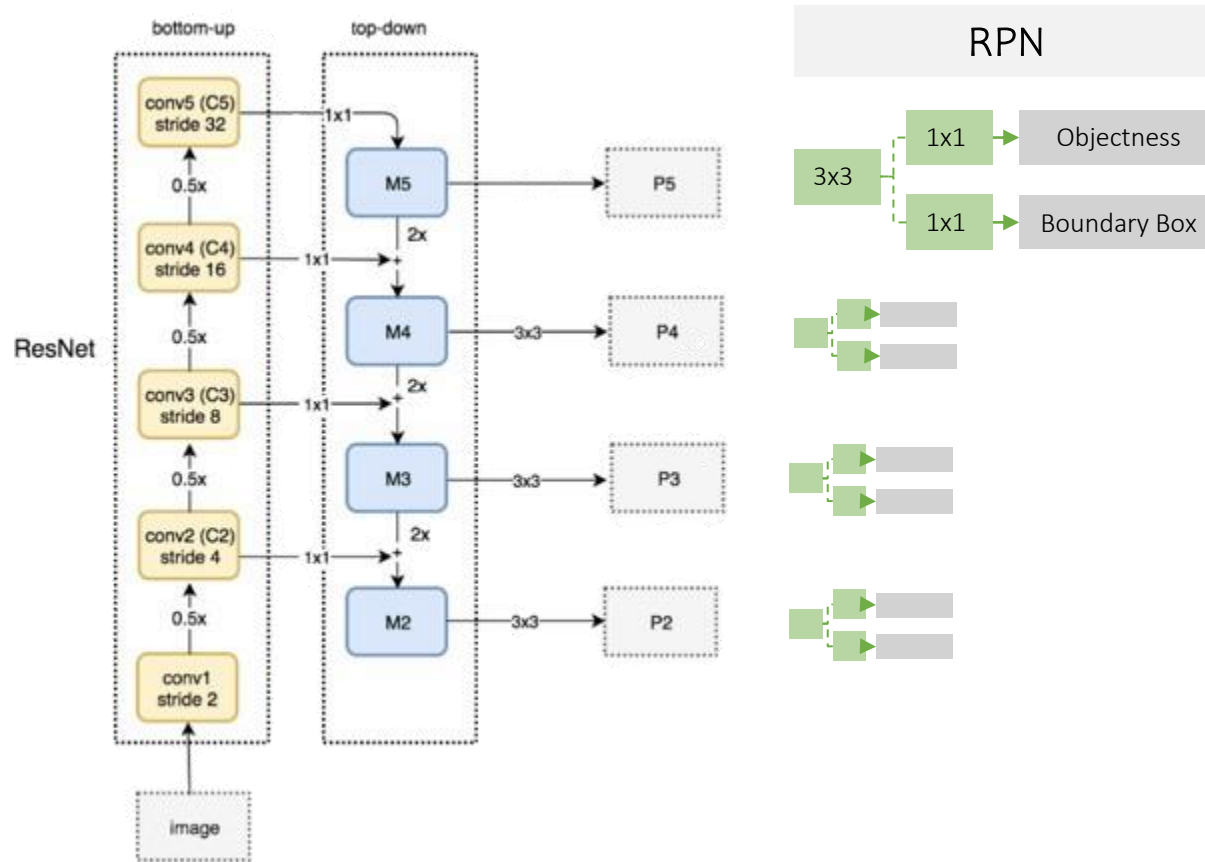
[2] [Networks in Networks and 1x1 Convolutions - Deep Convolutional Models: Case Studies | Coursera](#)



layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112			$7 \times 7, 64$, stride 2		
				3×3 max pool, stride 2		
conv2_x	56×56	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3_x	28×28	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
conv4_x	14×14	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	1×1		average pool, 1000-d fc, softmax			
FLOPs		1.8×10^9	3.6×10^9	3.8×10^9	7.6×10^9	11.3×10^9









FPN-based RPN generates more anchors because more than one feature maps are fed into RPN input.

First Stage of Faster R-CNN



Inference time for a frame : 3023 ms

RPN Out

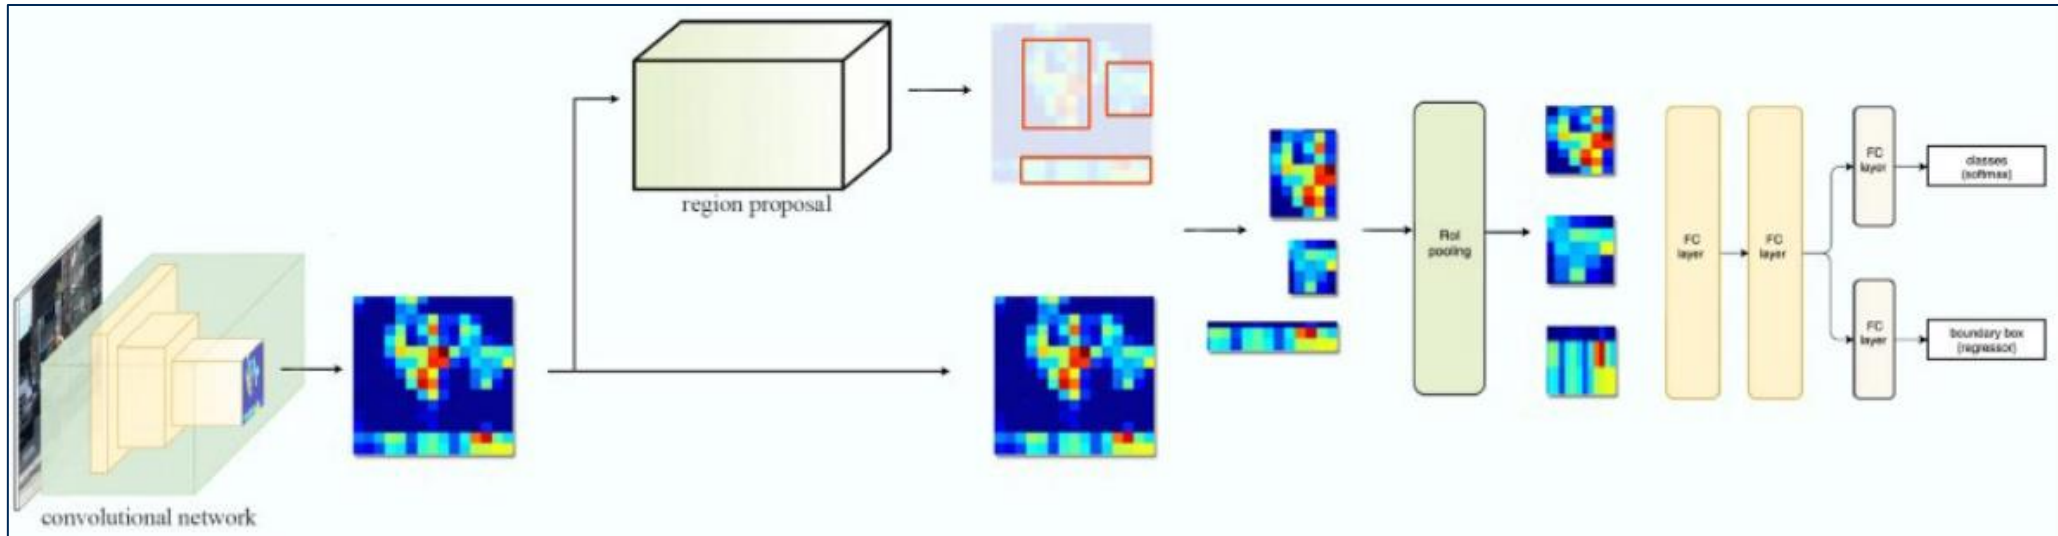


Inference time for a frame : 3926 ms

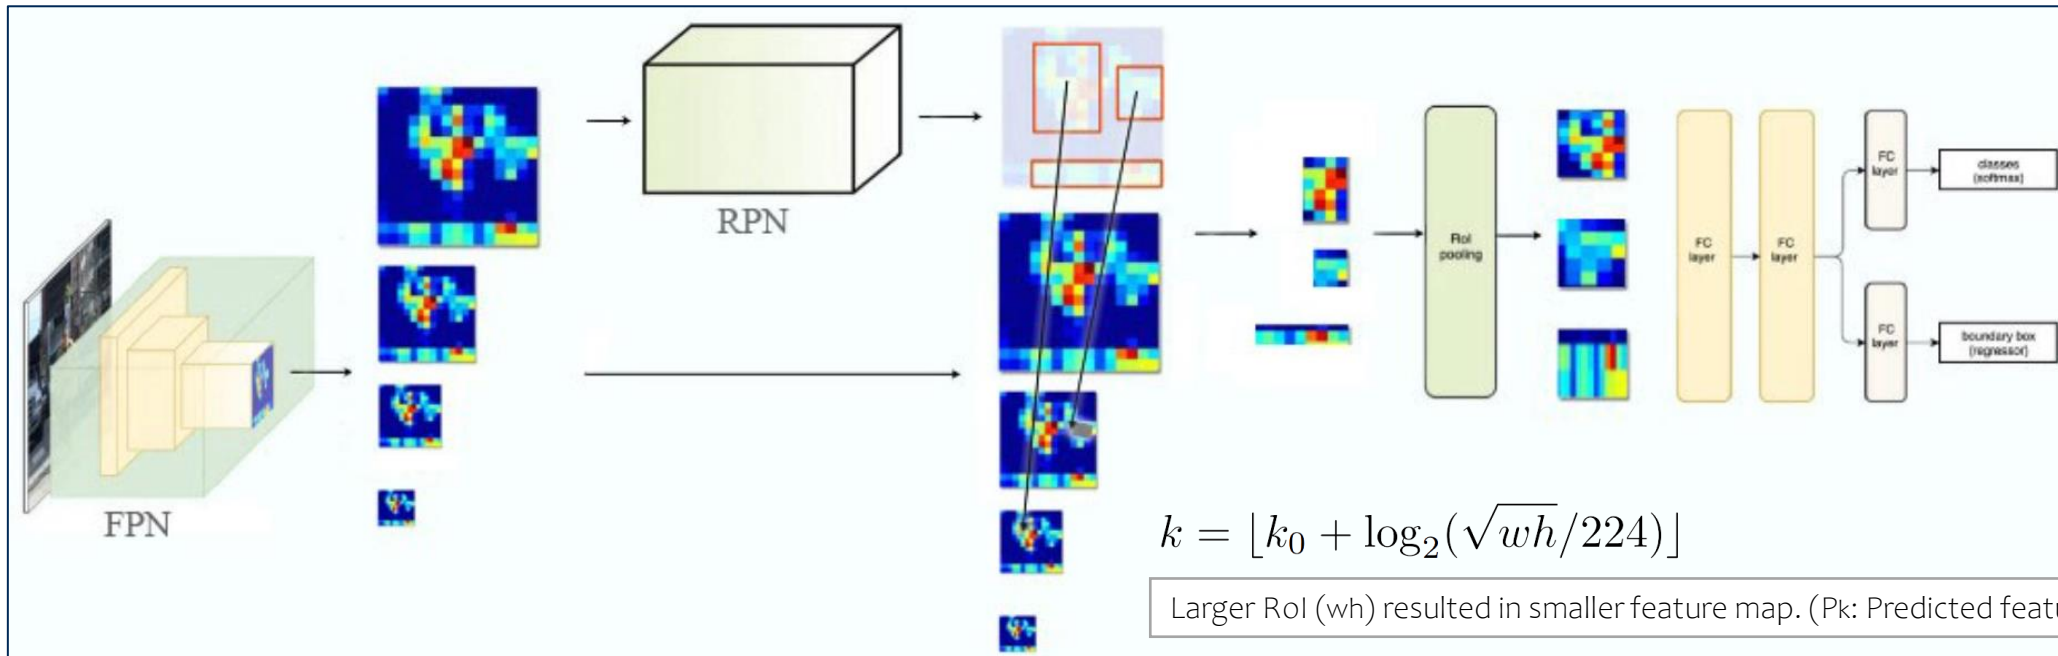
Final Detection Out



Faster R-CNN



FPN-based Faster R-CNN



RPN	feature	# anchors	lateral?	top-down?	AR ¹⁰⁰	AR ^{1k}	AR _s ^{1k}	AR _m ^{1k}	AR _l ^{1k}
(a) baseline on conv4	C_4	47k			36.1	48.3	32.0	58.7	62.2
(b) baseline on conv5	C_5	12k			36.3	44.9	25.3	55.5	64.2
(c) FPN	$\{P_k\}$	200k	✓	✓	44.0	56.3	44.9	63.4	66.2
<i>Ablation experiments follow:</i>									
(d) bottom-up pyramid	$\{P_k\}$	200k	✓		37.4	49.5	30.5	59.9	68.0
(e) top-down pyramid, w/o lateral	$\{P_k\}$	200k		✓	34.5	46.1	26.5	57.4	64.7
(f) only finest level	P_2	750k	✓	✓	38.4	51.3	35.1	59.7	67.6

FPN increases the #anchors generated in RPN even though scale variations are eliminated. (RPN anchor parameter k is less)

Faster R-CNN	proposals	feature	head	lateral?	top-down?	AP@0.5	AP	AP _s	AP _m	AP _l
(*) baseline from He <i>et al.</i> [16] [†]	RPN, C_4	C_4	conv5			47.3	26.3	-	-	-
(a) baseline on conv4	RPN, C_4	C_4	conv5			53.1	31.6	13.2	35.6	47.1
(b) baseline on conv5	RPN, C_5	C_5	2fc			51.7	28.0	9.6	31.9	43.1
(c) FPN	RPN, $\{P_k\}$	$\{P_k\}$	2fc	✓	✓	56.9	33.9	17.8	37.7	45.8

Inference Time on NVIDIA M40 GPU

0.148 sec

FPN-based Faster R-CNN (ResNet-50)

0.172 sec

Standard Faster R-CNN (ResNet-50)

FPN has extra layers but has a lighter weight head.

method	backbone	competition	image pyramid	test-dev					test-std				
				AP@.5	AP	AP _s	AP _m	AP _l	AP@.5	AP	AP _s	AP _m	AP _l
ours, Faster R-CNN on FPN	ResNet-101	-		59.1	36.2	18.2	39.0	48.2	58.5	35.8	17.5	38.7	47.8

Competition-winning *single-model* results follow:

G-RMI [†]	Inception-ResNet	2016		-	34.7	-	-	-	-	-	-	-	-
AttractionNet [‡] [10]	VGG16 + Wide ResNet [§]	2016	✓	53.4	35.7	15.6	38.0	52.7	52.9	35.3	14.7	37.6	51.9
Faster R-CNN +++ [16]	ResNet-101	2015	✓	55.7	34.9	15.6	38.7	50.9	-	-	-	-	-
Multipath [40] (on minival)	VGG-16	2015		49.6	31.5	-	-	-	-	-	-	-	-
ION [‡] [2]	VGG-16	2015		53.4	31.2	12.8	32.9	45.2	52.9	30.7	11.8	32.8	44.8

Thank you

